

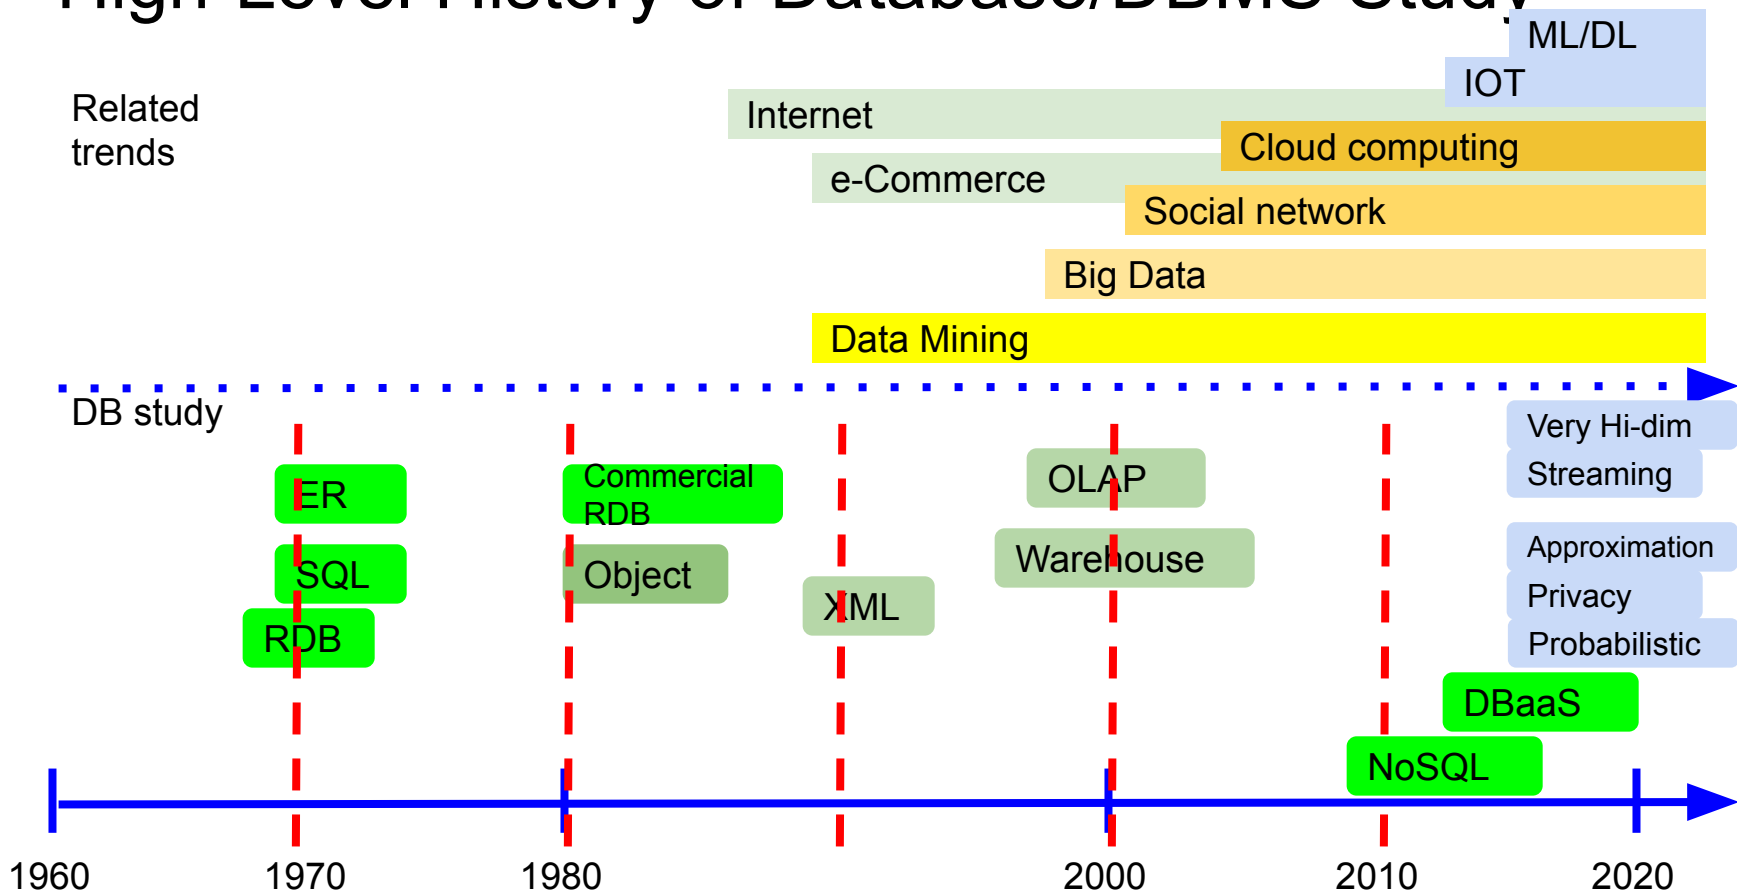
# DS108 Relational Database Overview

Ming-Ling Lo  
20191002

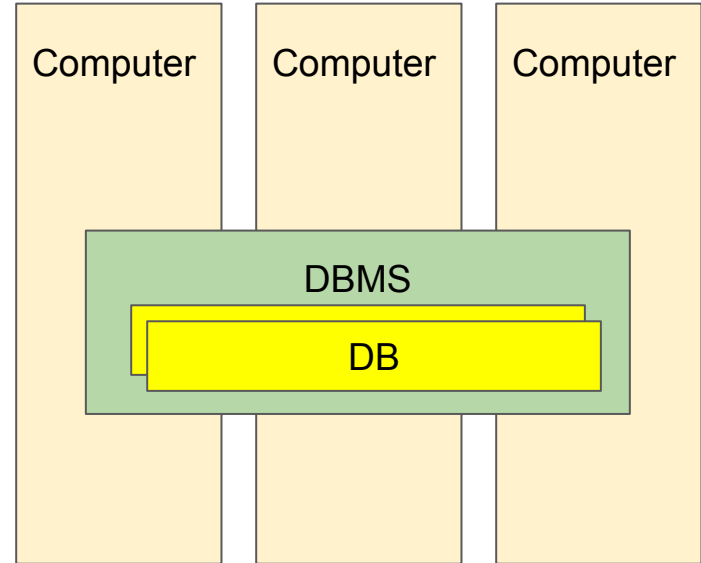
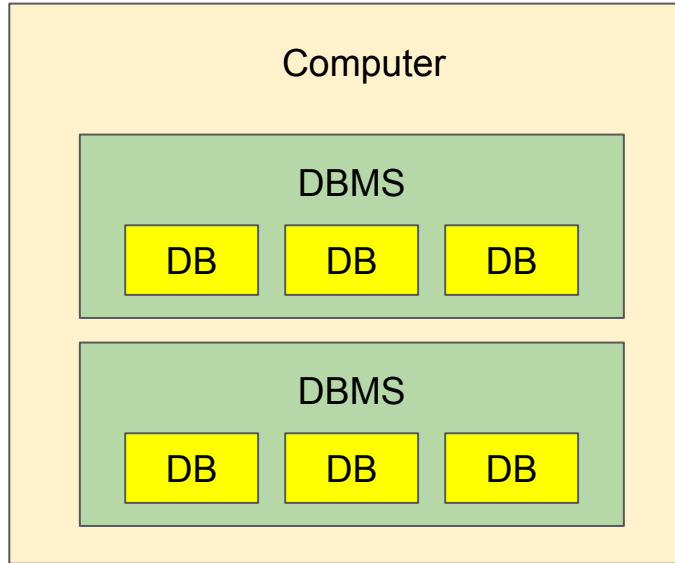
# The Study of Relational Databases

- Basics
  - Relational database model
  - SQL language
  - Relational Algebra
  - Relational Calculus
- Advanced basics
  - RDB design
    - Mapping from ER/EER to RDB, Normalization
  - Programming interfaces
    - ODBC, JDBC, Other programming
- Implementation and Performance

# High-Level History of Database/DBMS Study



# DB, DBMS Instances and Computers



Questions: Why storing data in the same DBMS into multiple databases?

# Relational Database Model

# Relational Database Model

- Used mathematical relation as its building block
- Theoretical basis in set theory and first-order predicate logic

<u>ID</u>	name
99633	John
99634	Mary
99635	Peter

=

- Intuitively
  - Think of Relation as a table
  - Row  $\rightarrow$  Tuple
  - Column  $\rightarrow$  Attribute
  - Data types of value  $\rightarrow$  Domain
  - But **Order of rows** is considered immaterial (not counted)

<u>ID</u>	name
99634	Mary
99633	John
99635	Peter

# Relational Database Formal Definition

- Three layers of definition
  - **Domain**
  - **Relation Schema; relation state**
  - **Database schema; database state**

# Values and NULLs in Tuples

- Atomic values

- Each value in a tuple must be an atomic value
- I.e., it is not divisible into components within the framework of the basic relational model
- Composite and multivalued attributes are not allowed.

- **NULL values**

- An important concept
- Used to represent the values of attributes that
  - Exist but we don't know
    - A student has a NULL for home phone (he doesn't have, or he has but we don't know.)
  - We don't know whether exist
    - A student has NULL for his office phones because he does not have office



# (Relational) Model Constraints

- Rules imposed on the modeling of data to make database safer, more meaningful and more useful
- Three main categories
  - [Inherent model-based constraints](#) or implicit constraints:
    - Constraints that are inherent in the data model
    - E.g. Tuples cannot repeat in a relation in RDB
  - [Schema-based constraints](#) or explicit constraints
    - Constraints that can be directly expressed in schemas of the data model, typically by specifying them in the DDL (data definition language)
  - [Application-based or semantic constraints](#) or business rules
    - Constraints that cannot be directly expressed in the schemas of the data model, and hence must be expressed and enforced by the application programs.

# (Relational) Model Constraints

- (Additional) rules that we impose on our modeling of data to make database safer, more meaningful and more useful
- Three main categories
- Inherent model-based constraints or implicit constraints:
  - Constraints that are inherent in the data model
  - E.g. Tuples cannot repeat in a relation in RDB
- Schema-based constraints or explicit constraints
  - Constraints that can be directly expressed in schemas of the data model, typically by specifying them in the DDL (data definition language)
- Application-based or semantic constraints or business rules
  - Constraints that cannot be directly expressed in the schemas of the data model, and hence must be expressed and enforced by the application programs.

# Key Constraints (1)

- A relation is defined as a set of tuples
- Since all elements of a set must be distinct; hence, all tuples in a relation must also be distinct
  - $\Rightarrow$  I.e., no two tuples can have the same combination of values for all their attributes.
- Usually, there may be **subsets of attributes** in a relation schema  $R$  with the property that no two tuples in any relation state  $r(R)$  should have the same combination of values for these attributes.
  - Denote one such **subset of attributes** by **SK**
  - $\Rightarrow$  For any two distinct tuples  $t_1$  and  $t_2$  in a relation state  $r$  of  $R$ , we have the constraint that:  
 $t_1[SK] \neq t_2[SK]$

# Key Constraints (2)

- Any such set of attributes SK is called a superkey of the relation schema R.
  - A superkey SK specifies a uniqueness constraint that no two distinct tuples in any state  $r(R)$  can have the same value for SK.
  - Every relation has **at least one default superkey** — the set of all its attributes
- Definition: a key K of a relation schema R
  - Is a superkey of R
  - Removing any attribute A from K, the rest of attribute set K' is no longer a superkey of R
- Properties of key
  - A key is also a superkey, but not vice versa.
  - Any superkey formed from one single attribute is always also a key

# Key Constraints (3)

- In general, a relation schema may have more than one key. In this case, each of the keys is called a candidate key.
- For example
  - A relation for CAR may have two candidate keys: License\_number and Engine\_serial\_number
  - A relation for student may have two candidate keys: student\_ID and National\_ID
- One of the candidate keys is designated as the primary key of the relation
  - The values of primary key are used to identify tuples in the relation.
  - Convention: the attributes that form the primary key of a relation schema are underlined
- Choice of one candidate keys to become primary key is somewhat arbitrary
  - Usually better to choose a primary key with a single attribute or a small number of attributes
  - The other candidate keys are designated as unique keys, and are not underlined

# Constraints on NULL


- Specifies whether NULL values are or are not permitted for a particular attribute
  - For example, if every STUDENT tuple must have a known, valid value for the Name attribute, then Name of STUDENT is constrained to be NOT NULL.

# Entity Integrity Constraints

- The entity integrity constraint states that no primary key value can be NULL.
  - Because the primary key value is used to identify individual tuples in a relation, having NULL values for the primary key implies that we cannot identify some tuples.
  - For example, if two or more tuples had NULL for their primary keys, we may not be able to distinguish them if we try to reference them from other relations.


<u>EID</u>	ENAME
101	John
102	Mary

<u>EID</u>	ENAME
101	John
NULL	Mary



# Referential Integrity Constraints

- The referential integrity constraint is specified between two relations
- is used to maintain the consistency among tuples in the two relations.
- Informally, the referential integrity constraint states that a tuple in one relation that refers to another relation must refer to an existing tuple in that relation.





# Transaction

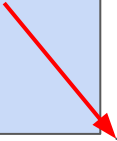
- A transaction
  - Is a unit of work that includes **several database operations**
    - such as reading from the database, or insertions, deletions, or updates to the database.
  - At the end of a transaction
    - It must leave the database in a valid state
    - Must satisfies all the constraints specified on the database schema.
    - All operations either all succeed together, or all failed together
      - Cannot let it happen  $\Rightarrow$  some succeed, some fail
    - These operations forms an **atomic unit of work** against the database.
  - For example, a transaction to apply a bank withdrawal will typically read the user account record, check if there is a sufficient balance, and then update the record by the withdrawal amount.

# SQL Overview

# Terminology about Database Sublanguages

Divide into details (informally, usually by commercial systems)

Broadly speaking

- DDL: data definition language
  - DML: data manipulation language
- 

- DDL: data definition language

- SDL: storage definition language

- DCL: data control language

- DQL: data query language
- DML: data manipulation language

# What is SQL?

- E.g.  
`SELECT * FROM department  
WHERE name = 'EE';`
- A language for defining your relational database model (DDL)
- A language to access relational data (DML)
- An Interface to Relational Database
- An important tool for DBMS
- A declarative language
- A language with “life”
- Borrowed from and extended to many applications beyond RDB
- We will present the “principles” of SQL
- SQL you encounter in actual systems in the future will all have small variations.

# SQL DDL

## (Data Definition Language)

# Create Table

- **CREATE TABLE PROJECT**  
( Pname VARCHAR(15) **NOT NULL**,  
Pnumber INT **NOT NULL**,  
Plocation VARCHAR(15),  
Dnum INT **NOT NULL**,  
**PRIMARY KEY** (Pnumber),  
**UNIQUE** (Pname),  
**FOREIGN KEY** (Dnum)  
**REFERENCES** DEPARTMENT(Dnumber));

# Specifying Key Constraints

- **PRIMARY KEY** clause: specify primary key of a relation
  - If a primary key has a single attribute, the clause can follow the attribute directly.
  - For example, the primary key of PROJECT can be specified as follows Dnumber INT PRIMARY KEY;
- **UNIQUE** clause: specify secondary key of a relation
  - The UNIQUE clause can also be specified directly for a secondary key if the secondary key is a single attribute,
  - For example: Pname VARCHAR(15) UNIQUE;

# Specifying Key Constraints (2)

- **Example**

```
CREATE TABLE PROJECT  
( Pname VARCHAR(15) NOT NULL,  
  Pnumber INT NOT NULL,  
  Plocation VARCHAR(15) DEFAULT 'TAIPEI',  
  Dnum INT NOT NULL CHECK (Dnum > 0  
    AND Dnum <21),  
  PRIMARY KEY (Pnumber),  
  UNIQUE (Pname),  
  FOREIGN KEY (Dnum)  
  REFERENCES DEPARTMENT(Dnumber));
```

- **Example**

```
CREATE TABLE PROJECT  
( Pname VARCHAR(15) UNIQUE NOT  
  NULL,  
  Pnumber INT PRIMARY KEY NOT NULL,  
  Plocation VARCHAR(15) DEFAULT 'TAIPEI',  
  Dnum INT NOT NULL CHECK (Dnum > 0  
    AND Dnum <21),  
  FOREIGN KEY (Dnum)  
  REFERENCES DEPARTMENT(Dnumber));
```



# Specifying Referential Integrity Constraints

- FOREIGN KEY clause

- **Example**

**CREATE TABLE** PROJECT

( Pname VARCHAR(15) **NOT NULL**,

Pnumber INT **NOT NULL**,

Plocation VARCHAR(15) DEFAULT 'TAIPEI',

Dnum INT **NOT NULL CHECK** (Dnum > 0 AND Dnum <21),

**PRIMARY KEY** (Pnumber),

**UNIQUE** (Pname),

**FOREIGN KEY** (Dnum)

**REFERENCES** DEPARTMENT(Dnumber));

What does this constraint really regulate?

In what situation will this command fail?

Does it affect other commands?

# Schema Changing: Drop

- Drop
  - Can drop database, schema, catalog
  - Can drop table
  - Restrict or cascade clause: decide what to do with affected objects
- DROP DATABASE university (mySQL)
- DROP SCHEMA university
  - In MySQL, you can say DROP DATABASE ...
- DROP SCHEMA university CASCADE
- DROP TABLE departments RESTRICT
- DROP TABLE departments CASCADE

# Schema Changing: Alter

- ALTER {DATABASE|SCHEMA} db\_name alter\_specification ...;
- ALTER TABLE: change the details of a table definition
  - Can add, drop column, constraint, index
  - Can rename column or change the definition of column
- ALTER TABLE students **ADD COLUMN** grade DECIMAL(5.2);
- ALTER TABLE students **DROP COLUMN** grade CASCADE;
  - We can specify CASCADE or RESTRICT when drop column
- ALTER TABLE students ADD COLUMN grade int DEFAULT 1.0;
- ALTER TABLE students **ALTER COLUMN** grade SET DEFAULT 2.0;
- ALTER TABLE students ALTER COLUMN grade DROP DEFAULT;


# SQL DML

## (Data Manipulation Language)

# Insert

- Three kinds of usage
  - Specific all values
  - Specific attribute names and values
  - Get value from a select clause
- Create table students (  
ID int, name char(20), department char(20));
- **Insert into students values (1, “John smith”, “Math”);**
- **Insert into students (ID, name) values (2, “Peter Chen”);**
- **Insert into students**  
**Select HID, Hname, Wishdept From highschool\_students**  
**Where status = ‘selected’;**

# Delete

- Two types of usage
  - With conditions
  - Without condition
- Create table students (  
ID int, name char(20), department char(20));
- **Delete from students**  
**Where department = 'math';**  Tuple select condition
- **Delete from students;**
  - Delete all rows in the table 'students' ⇒ Be careful
  - But the schema of 'students' still remains

# Update

- Modify attribute values of one or more selected tuples
  - Specify what values to update, and which tuple to update to
- Update tablename set attribute='value', attribute='value'... where condition
- Can also set to NULL and set to default

What values to update

Which tuple to update to

- Examples

- UPDATE PROJECT SET Plocation = 'Taipei', Dnum = 5 WHERE Pnumber = 10;
- UPDATE EMPLOYEE SET Salary = Salary \* 1.1 WHERE Dno = 5;
- UPDATE EMPLOYEE SET Salary = DEFAULT WHERE Dno = 5;
- UPDATE EMPLOYEE SET Salary = NULL WHERE Dno = 5;

# Select (Basic)

- Most important command in SQL
- The way to execute query and get answers!
- Basic format

**SELECT <attribute list>**

**FROM <table list>**

**WHERE <condition>;**

- Note: when there are multiple tables in the <table list>, it means some kind of join operation
- <condition> typically contains atomic condition expressions connected by AND, OR, NOT



## Select (2)

- Create table students (  
ID int, name char(20), deptID int);
- Create table departments (  
DID int, Dname char(20));
- Select ID, name  
From students  
Where deptID = 3;
- Select ID, name  
From students, departments  
Where Dname='EE' and deptID=did;

## Select (2-1)

- Select ID, name  
From students, departments  
Where Dname='EE' and DEPTID = DID;

Student: Referencing relation

<u>ID</u>	NAME	DEPTID
101	John	7
203	Mary	8
307	Tom	NULL

Department: Referenced relation

<u>DID</u>	DNAME	
7	EE	
8	CS	
15	Video game	

Referential integrity holds between Student and Department now

# Transaction

- Basic syntax:
  - START TRANSACTION;  
    <sql commands>...  
    <sql commands>...  
    <sql commands>...  
    COMMIT (or ROLLBACK)
  - SET autocommit = {0|1}
    - Default is '1'
- START TRANSACTION;  
SELECT @A:=SUM(salary) FROM table1 WHERE type=1;  
UPDATE table2 SET summary=@A WHERE type=1;  
COMMIT;